

# SC1015 Intro to DSAI:

Enhancing road safety in California with data-driven insights

Joanna Wu Haoyue [U2322092E]  
Sricharan Balasubramanian [U2322339L]  
Travis Tan Hai Shuo [U2322928D]

Lab REP2 Team 4

# Chosen Dataset

<https://www.kaggle.com/datasets/sobhanmoosavi/us-accidents>

## US Accidents (2016 - 2023)

A Countrywide Traffic Accident Dataset (2016 - 2023)

BY: SOBHAN MOOSAVI



Data Card Code (416) Discussion (46) Suggestions (0)

### About Dataset

#### Description

This is a countrywide car accident dataset that covers **49 states of the USA**. The accident data were collected from **February 2016 to March 2023**, using multiple APIs that provide streaming traffic incident (or event) data. These APIs broadcast traffic data captured by various entities, including the US and state departments of transportation, law enforcement agencies, traffic cameras, and traffic sensors within the road networks. The dataset currently contains approximately **7.7 million** accident records. For more information about this dataset, please visit [here](#).

#### Usability ⓘ

10.00

#### License

[CC BY-NC-SA 4.0](#)

#### Expected update frequency

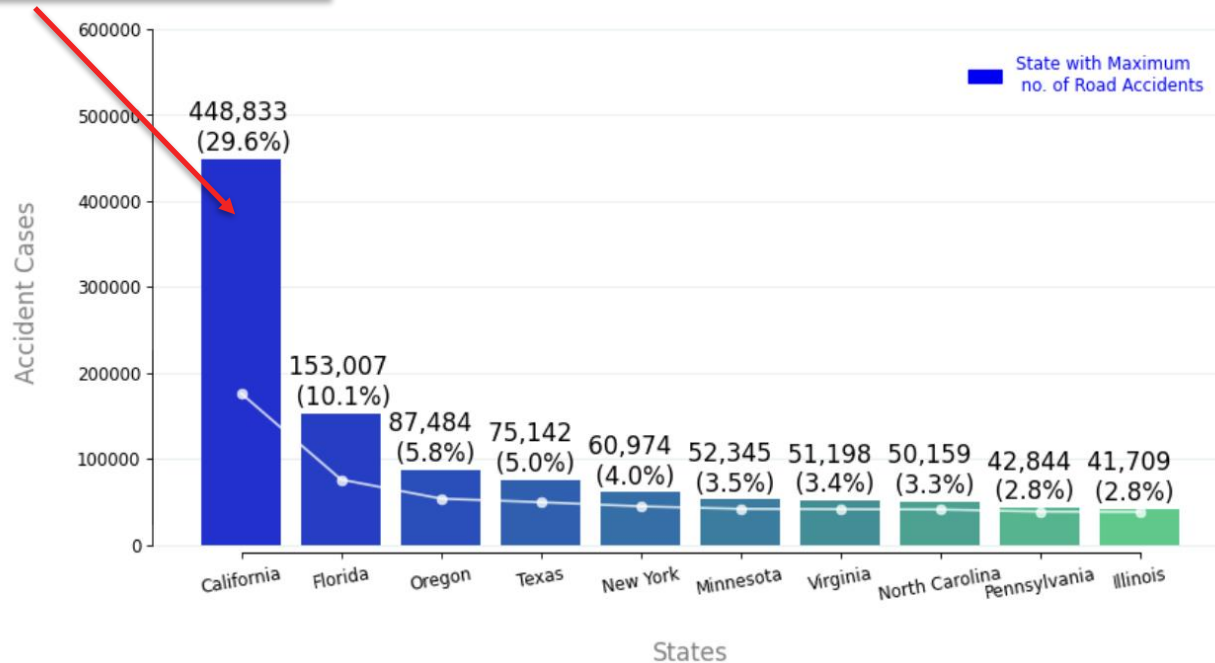
Never

#### Tags

# Chosen Dataset

We will be heading to California for exchange in August!

Top 10 States with most no. of Accident cases in US (2016-2020)





# Problem Definition

Analyse the **times of day/night** and **weather conditions** where Californian **cities** are **most dangerous to drive in**.

1. This helps emergency services **allocate resources optimally**, **reducing fatality of road accidents**.
2. This can be integrated into GPS to **caution drivers** when they are driving in accident-prone conditions, **reducing frequency of road accidents**.



# Data Preparation & Cleaning



## Data Type Conversion

Street	object
City	object
County	object
State	object
Zipcode	object



Street	category
City	category
County	category
State	category
Zipcode	category

Dataset's features were **converted into appropriate types**, making data handling easier for ML models.

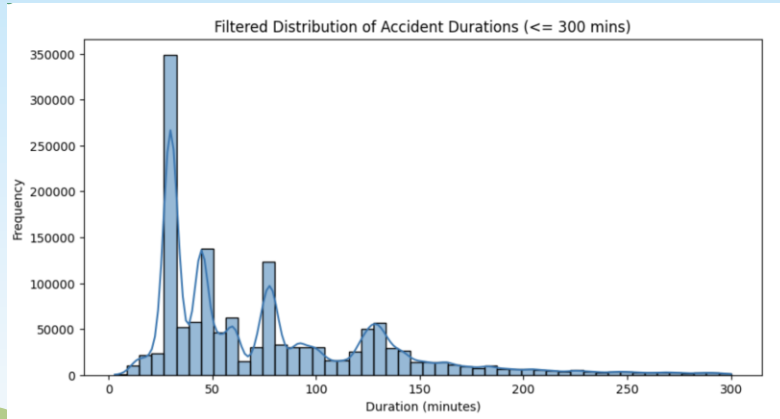
```
df['Start_Time'] = pd.to_datetime(df['Start_Time'], errors='coerce')  
df['End_Time'] = pd.to_datetime(df['End_Time'], errors='coerce')
```

Time elements were **parsed** (time columns were converted into datetime objects).

Date and time were separated. Only **time feature was extracted** as date was meaningless in the problem context.

# Feature Engineering

## Feature Creation



```
df['Duration'] = df['End_Time'] - df['Start_Time']
```

Created a new feature 'Duration' to visualise duration of accidents, and to help decide if 'End\_Time' is relevant to our analysis

## Feature Selection

Number of unique...

State	1
Zipcode	129022
Country	1
Timezone	2


We dropped features that were irrelevant to our problem statement. E.g. 'Country', 'State' etc. are the same for each data point, and thus can be dropped. Twilights are alternative ways to measure time. 'End\_Time' is shown to be irrelevant

# Handling Missing Data



## Critical features with low NULL %

These rows were **dropped** as the change in results is **negligible**, but quality of dataset will be **greatly improved**.



## Weather features

These cells were filled with **median values** to reflect the central tendency of weather conditions



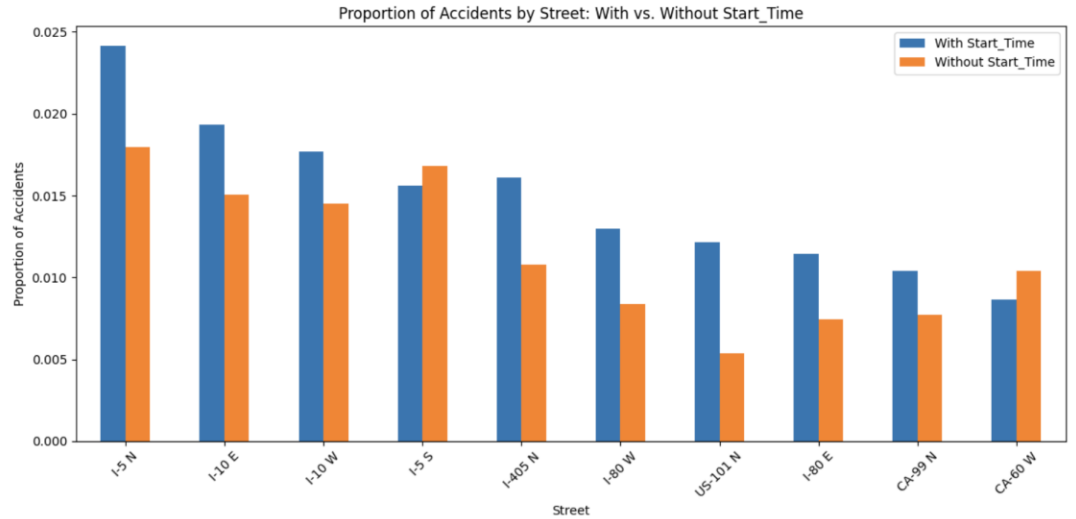
## High NULL %

Some features (Wind\_Chill) were dropped as it is **strongly correlated** with other features like Temperature. **Precipitation NULLs** are inputted with 0.

# Handling Missing Data

'Start\_Time' → 10% of values are NULL!

To decide how to handle NULL accident time values, we compared the **proportions** of the most accident-prone streets that have a recorded time and those that do not. Ultimately, **no major disproportion** was observed, implying the results would not be heavily skewed with the removal of these NULL points.





# Outlier Handling



## Data Normalisation

Noticing a wide spread of numerical data (many unique values), we adjusted the data so the values fall within a **specific range**



## Outlier Removal

Outliers were detected in various numerical categories upon observing the spread of data in boxplots. They were removed using the **Interquartile** method.



# Exploratory Data Analysis

## Regions

Are trends **homogenous** all across California, or are there **specific places** to consider?

## Day and Time

What trends can we identify in the **frequency** of accidents against the **time of the day**, and the **day of the week**?

## Weather Conditions

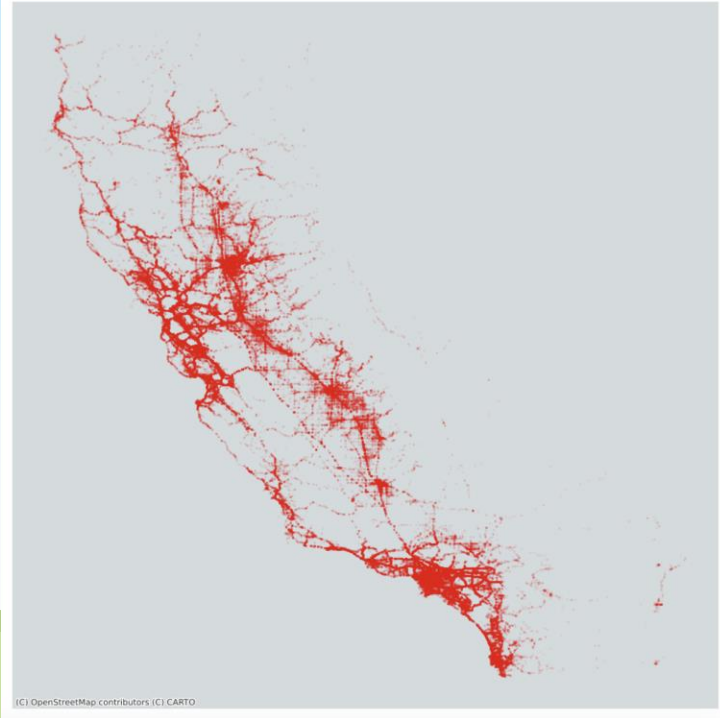
In the regions, do accidents occur in all weather conditions or only in **certain weather conditions**?

# General Regional Analysis

We see that the distribution is centred in some specific parts of California, rather than being distributed evenly all around the state.

Narrows scope of our discussion and analysis

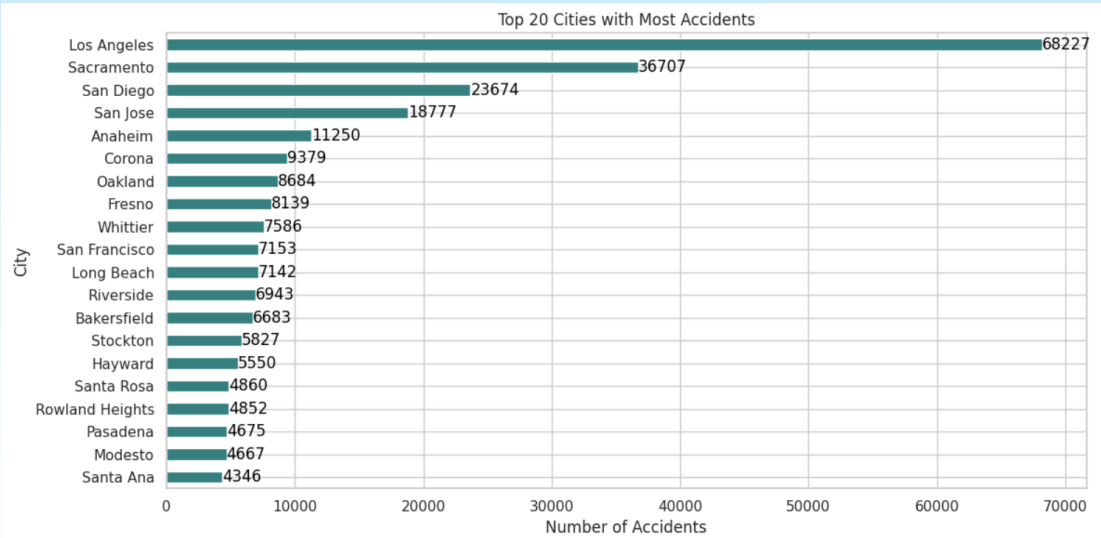
Accident Density Map: California



# City Analysis


We see a clear spike in some cities, and this validates our approach of looking at cities within California, rather than treating the entire state as one big homogenous space

Care must be taken within cities, and this is a foundation for our future model




# Street Analysis

The streets show a more homogenous mix, but the streets remain meaningful for us to look at due to their direct relevance to our use case. However, the cardinality is a weakness that makes it hard to be analysed




# Time Analysis

We see a local spike between 7AM and 9AM, and a global spike between 4PM and 7PM. These are peak hours, and clearly this is another dimension for our analysis




# Day Analysis

A clear distinction between weekdays and weekends. A specific spike on Fridays, possibly an indicator of more haphazard driving



# Cross Validation - City

2 clear spikes seen across all cities during the peak hours in the morning and evening









# Cross Validation - City

The day trends are relevant to the cities with largest number of accidents, validating our focus on it



# Cross Validation - City


There are no major outliers to this day-based trend. So, our analysis can be used to explain accident trends across the state without concerns of over-generalisation of trends without a regional focus.



# Weather Analysis

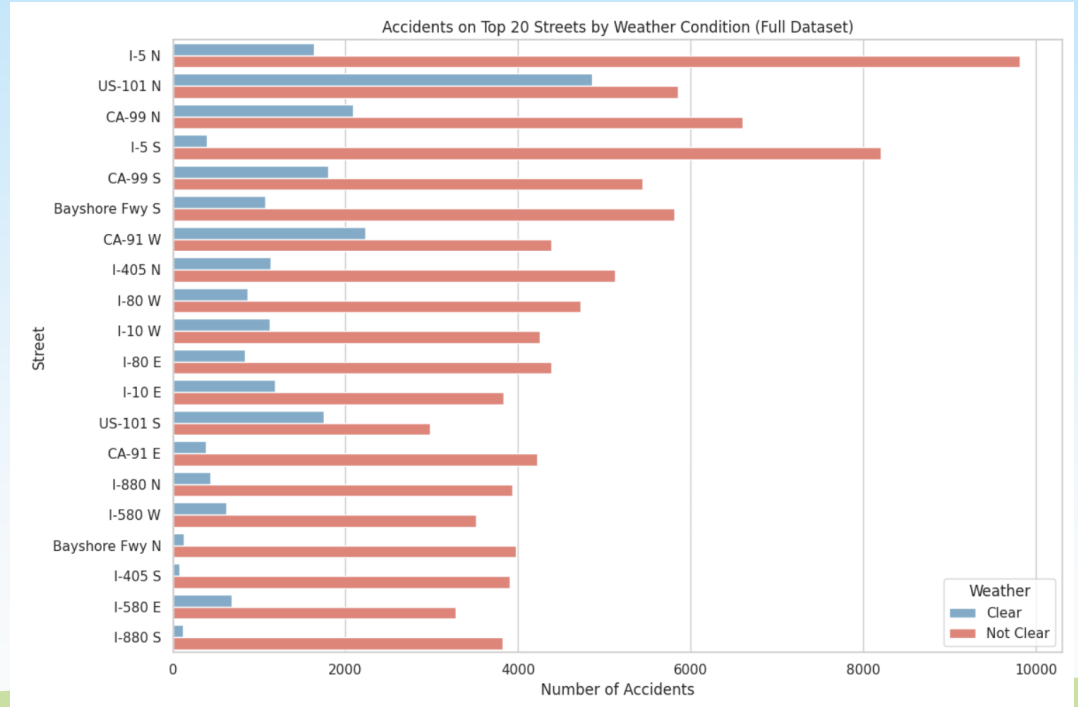
Unable to remove any dimensions for weather metrics, due to low cross-correlations. All these factors remain in the model.

Descriptions have high cardinality, bucketed into 6 binary flags - 'Is\_Windy', 'Is\_Stormy', 'Is\_Rainy', 'Is\_Foggy', 'Is\_Snowy', 'Is\_Clear'



# Weather Analysis

Generally we see that unclear weather contributes heavily to the occurrence of accidents, making it a key factor that we wish to consider. The different factors of weather can be analysed to make predictions.





# EDA Conclusions

## Regions

Base core analysis on the cities and streets (could have issues), as they **show trends** and also are the most **actionable**

## Day and Time

Do analysis based on the **peak hours** of the day, and whether it is a **weekday or weekend**

## Weather Conditions

Consider all **non-clear weather** conditions, **less** so for specific **roads with a homogeneous split**

# Machine Learning

There were 3 approaches for our ML algorithm:

Accident Classification  
(Classification)

Not possible as we possess a **positive-only dataset**, leading us to be unable to distinguish between accidents and non-accidents



Rate of Accidents Prediction  
(Regression)

Possible use of linear regression to determine the rate of accidents in a city over the next hour based on additional information such as time, weather and location



Severity of Accident  
(Classification)

Extremely low quantities of high severity data, resulting in imbalance in the data set



# Reducing Dataframe Size

Based on the EDA conducted, we have identified a few key metrics that could affect the accident rates. From factors such as the **City, Day of the Week as well as Weather condition**. These variables will be extracted to trim the dataset for faster processing by the algorithms. Street is removed here due to its extremely high cardinality, instead focusing on City.

```
columns_to_keep = [  
    'Start_Time', # Time information  
    'City', #City  
    'Temperature(F)', # Weather data  
    'Humidity(%)', # Weather data  
    'Pressure(in)', #Pressure data  
    'Wind_Speed(mph)', # Weather data  
    'Start_Hour', # Extracted from Start_Time  
    'Day_of_Week', # Extracted from Start_Time  
    'Is_Windy',  
    'Is_Stormy',  
    'Is_Rainy',  
    'Is_Foggy',  
    'Is_Snowy',  
    'Is_Clear'  
]
```



# Further Feature Engineering

## Feature Creation

```
# Step 1: Floor to hour
df_algo1['Start_Hourly'] = df_algo1['Start_Time'].dt.floor('H')

# Step 2: Create "Next Hour"
df_algo1['Next_Hourly'] = df_algo1['Start_Hourly'] + pd.Timedelta(hours=1)

# Step 3: Count future accidents per city per hour
city_hourly_counts = df_algo1.groupby(['City', 'Start_Hourly']).size().rename('Accident_Count')

# Step 4: Shift within each city to get "next hour's" count
city_hourly_next = city_hourly_counts.groupby(level=0).shift(-1).rename('Accident_Next_Hour')

# Step 5: Combine the City/Hour index with the shifted counts
target_df = pd.concat([city_hourly_counts, city_hourly_next], axis=1).reset_index()

# Step 6: Merge back into main df
df_algo1 = df_algo1.merge(target_df[['City', 'Start_Hourly', 'Accident_Next_Hour']],
                        on=['City', 'Start_Hourly'], how='left')

print(df_algo1.head(3))
```

Created a new feature 'Accident\_Next\_Hour' based on the city and hour as an indicator of the accidents occurring in the next hour

## Target Variable

Accident_Next_Hour
0.0
1.0
0.0
0.0
0.0

This new feature would act as the target variable for our models to predict

## Further Feature Engineering

# Feature Creation

```
# Step 1: Make sure your time is rounded to the hour
df_algol['Start_Hourly'] = df_algol['Start_Time'].dt.floor('H')

# Step 2: Count number of accidents per city per hour
city_hourly_counts = df_algol.groupby(['City', 'Start_Hourly']).size().rename('Accident_Current_Hour')

# Step 3: Create the Lag (previous hour's accident count)
city_hourly_counts_lag = city_hourly_counts.groupby(level=0).shift(1).rename('Accident_Prev_Hour')

# Step 4: Combine the counts into a DataFrame for merging
lag_features = pd.concat([city_hourly_counts, city_hourly_counts_lag], axis=1).reset_index()

# Step 5: Merge Lag features into main DataFrame
df_algol = df_algol.merge(lag_features, how='left', on=['City', 'Start_Hourly'])
```

```

Accident_Prev_Hour    0.0
                      0.0
                      0.0
                      0.0
                      0.0
                      0.0
                      0.0
                      0.0
                      0.0
                      0.0
                      0.0
                      0.0
                      0.0
                      0.0
                      1.0
                      0.0
                      0.0
                      0.0
                      0.0

```

Created a new feature 'Accident\_Prev\_Hour' based on the city and hour as an indicator of the accidents occuring in the previous hour, to feed time lag data into our model

# Feature Creation

```
# Hour of the day (0 to 23)
df_algo1['Start_Hour'] = df_algo1['Start_Time'].dt.hour

# Convert hour into sine and cosine to capture cyclic nature
df_algo1['Hour_Sin'] = np.sin(2 * np.pi * df_algo1['Start_Hour'] / 24)
df_algo1['Hour_Cos'] = np.cos(2 * np.pi * df_algo1['Start_Hour'] / 24)

df_algo1.head()
```

Hour_Sin	Hour_Cos
0.500000	-0.866025
0.500000	-0.866025
0.500000	-0.866025
0.258819	-0.965926
0.258819	-0.965926

Created new features 'Hour\_Sin' and 'Hour\_Cos' to better reflect the cyclic nature of the hours of the day

# Data Transformation

## Target Encoding

```
# Calculate the mean of the target variable for each city
city_target_mean = df_algo1.groupby('City')['Accident_Next_Hour'].mean()

# Map the mean target value to the 'City' column
df_algo1['City_Encoded'] = df_algo1['City'].map(city_target_mean)

# Check the result
print(df_algo1.head(20))
```

```
City_Encoded \
0.227317
0.749108
0.119518
0.022876
0.132522
0.139405
1.843790
0.139405
0.396246
0.013158
0.108696
0.147407
0.147407
0.025370
0.010135
0.108696
0.034296
0.022000
0.340076
0.156190
```

Due to the **high cardinality** of the City variable, One Hot Encoding was not possible, while Label Encoding could introduce unintended patterns, use **Target Encoding**

## Ordinal Encoding

```
day_of_week_map = {
    'Monday': 0,
    'Tuesday': 1,
    'Wednesday': 2,
    'Thursday': 3,
    'Friday': 4,
    'Saturday': 5,
    'Sunday': 6
}

# Apply the mapping to the 'Day_of_Week' column
df_algo1['Day_of_Week'] = df_algo1['Day_of_Week'].map(day_of_week_map)

# Check the result
print(df_algo1[['Day_of_Week']].head())
```

Since Day\_Of\_Week has an inherent order, it can be converted by undergoing a simple **Ordinal Encoding**

# Hyperparameter Tuning

```
from sklearn.ensemble import RandomForestRegressor

# Define the hyperparameters for Random Forest
rf_param_grid = {
    'n_estimators': [100, 200, 300],
    'max_depth': [None, 10, 20, 30],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'bootstrap': [True, False]
}
```

Random Forest

```
from sklearn.ensemble import GradientBoostingRegressor

# Define the hyperparameters for Gradient Boosting
gb_param_grid = {
    'n_estimators': [100, 200],
    'learning_rate': [0.01, 0.1, 0.2],
    'max_depth': [3, 5, 7],
    'subsample': [0.8, 1.0],
    'min_samples_split': [2, 5],
    'min_samples_leaf': [1, 2]
}
```

Gradient Boosting

```
from sklearn.linear_model import Ridge

# Define the hyperparameters for Ridge
lr_param_grid = {
    'alpha': [0.1, 1, 10, 100]
}
```

Linear Regression

# Model Results

	MSE	MAE	$R^2$
Random Forest	1.6470	0.6364	0.4834
Gradient Boosting	1.7667	0.6831	0.4458
Linear Regression	1.9202	0.7316	0.3977

From our results, it is clear that Random Forest Regressor performed the best with the lowest MSE and MAE and highest  $R^2$



## Model Results

But why is the  $R^2$  value of the model still so low?


From our results, it is clear that Random Forest Regressor performed the best with the lowest MSE and MAE and highest  $R^2$

# Model Results

Looking at the bins of possible accidents in the next hour, it is clear that the majority lie in the range of 0-2. As a result, the model could be overfitting as it cannot account for the outliers in the regression. Could converting this into a classification problem be better?

```
df_algol['Accident_Next_Hour'].hist(bins=30)
```

<Axes: >



# Model Results

Using a different approach, we now classify the rate of accidents into 4 buckets, none, low, medium and high.

```
def categorize_accidents(x):  
    if x == 0:  
        return 'none'  
    elif x <= 2:  
        return 'low'  
    elif x <= 5:  
        return 'medium'  
    else:  
        return 'high'  
  
df_algo1['Accident_Class'] = df_algo1['Accident_Next_Hour'].apply(categorize_accidents)  
df_algo1.head()
```



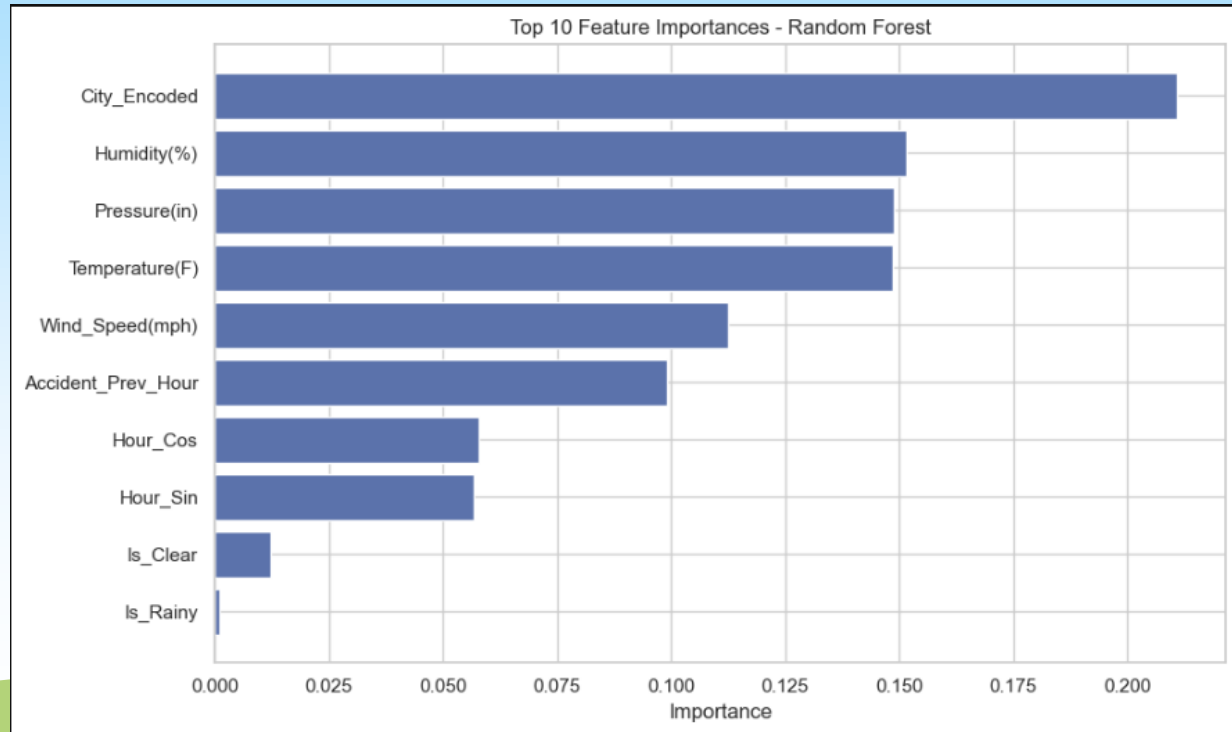
# Model Results

Out of Random Forest, Gradient Boosting, Logistic Regression, XGBoost and HistGradientBoosting, **Random Forest performed the best**. It could quite accurately predict no accidents (3) with F1 score of 0.93 and high rates of accidents (0) with F1 score of 0.80, but struggled with predicting in between, at low (2) and especially as medium (1) rates of accidents.



# Feature Importance

Looking at the feature importance for the Random Forest, it is clear that City is the greatest predictor of the rate of accident.



# Conclusion

## Problem Addressing

Unable to predict numeric probability of accidents due to limitations (lack of negative data), but identified key determinants to look at

## Interesting Insights

Oddly, humidity seems to have the largest impact on the rate of accidents other than the city we are looking at, while conditions like pressure also seem to have large impacts

## Use Case and Future Scope

Use in our own road trips in California

Import traffic details for live inputs and negative data as we drive

# Thank you!

Joanna Wu Haoyue [U2322092E]  
Sricharan Balasubramanian [U2322339L]  
Travis Tan Hai Shuo [U2322928D]

Lab REP2 Team 4